

**[ PLACEHOLDER ]**

LLM Inference Request Lifecycle: Prefill → Decode + KV Cache

Export from `plot/llm_inference_lifecycle.drawio` and overwrite this file