

LLM Inference Request Lifecycle: Prefill → Decode, bound to the KV Cache (background)

